Claim summary: precision-recall frontier, not a single universal winner

Method / lambda	validation/RMSD role	precision role	recall role	final interpretation
harmonic prior	Strong global-prior baseline; weaker than hybrid lambda=5.0 on optimized loss/RMSD.	AMR-P weaker than all tested hybrid lambdas in the lambda comparison.	Best AMR-R within the harmonic-vs-hybrid lambda comparison.	Recall anchor for the harmonic/hybrid family; not dominated by hybrid on ensemble metrics.
hybrid lambda=0.025	Low-modulation hybrid; not selected by RMSD/validation loss.	Small AMR-P improvement over harmonic, far weaker than lambda=5.0.	Best hybrid COV-R and AMR-R; balanced by rank aggregation.	Best hybrid recall/balanced setting, but not a full dominance result over harmonic.
hybrid lambda=5.0	Strongest optimized validation loss and single-sample RMSD among harmonic/hybrid settings.	Precision-optimal by AMR-P; COV-P tied at 1.0.	Worse COV-R and AMR-R than harmonic and low-lambda hybrids.	RMSD/precision-favorable setting; not ensemble-optimal.
stiffness structured	Better than isotropic/torsion-only on tracked RMSD, weaker than harmonic/hybrid.	Improves AMR-P over harmonic but not as strongly as hybrid lambda=5.0.	Does not beat harmonic on AMR-R in the all-method COV/AMR summary.	Local stiffness alone helps precision but does not replace harmonic/global prior.
isotropic	Weak optimized baseline on loss/RMSD.	Poor AMR-P relative to harmonic and hybrids.	Not recall-leading in the final comparison.	Naive Cartesian source baseline remains weaker.
torsion only	Very weak optimized loss/RMSD.	Poor AMR-P and lower COV-P, indicating weak precision.	Recall-biased COV-R/AMR-R behavior, but not balanced.	Diversity/recall-favorable but precision-poor; useful as a tradeoff contrast.